

Can Blood Vessel Networks in Ultrafast Breast MRI Predict Breast Cancer Treatment Response?

Spencer Venancio¹ Sarit Bose² Aakrithi Ram³ Advisor: Anna Woodard, PhD⁴

¹University of Wisconsin–Madison

²Neuqua Valley High School

³Bergen County Academies

⁴Data Science Institute, University of Chicago

Why Might Blood Vessels Predict Treatment Response?

Neoadjuvant treatment (NAT) varies by breast cancer subtype and may include chemotherapy and HER2-targeted therapy. Only some patients achieve a **pathologic complete response (pCR)**, which is associated with improved long-term outcomes [1]. Earlier prediction could avoid ineffective treatment for women unlikely to benefit and identify nonresponders who need intensified or alternative therapy.

Tumors recruit abnormal vessels that can be dense, tortuous, leaky, and disorganized [3]. Their structure and contrast-enhancement dynamics may reflect tumor biology and affect drug delivery. We test whether these vessel-network features improve pCR prediction beyond simple summaries of the same MRI measurements.

Research Question: does the structure of breast blood vessel networks and the way contrast changes through them improve response prediction beyond simple summaries of the same MRI measurements?

UChicago Cohort: 179 Labeled, 1,226 Unlabeled Exams

UChicago contributed **181 labeled exams from 143 patients; 179 exams from 141 patients** are analyzed here (115 non-pCR / 64 pCR, 36%). Two exams lacked the required centerline, support mask, and aligned MRI.

Acquisition timing is heterogeneous: exams contain 13–24 frames spanning 96–480 s. Longitudinal exams are included. Cross-validation is grouped by patient, so one patient's exams never cross training/evaluation splits.

Pretraining draws on a separate, non-overlapping pool: **1,226 unlabeled exams from 225 patients** (1,081 CAPS, 145 CHIMEC), with no patient or exam shared with the labeled cohort.

Three Graph Representations Test Whether Topology Matters




Figure 1. The same vessel centerline represented three ways.

We represent the vascular network of each patient's breasts as a graph:

- **Voxel graph:** each vessel-centerline voxel is a node; neighboring centerline voxels are connected.
- **Junction graph:** branch points are nodes ($v \in V_{\text{vox}}, \deg(v) > 2$); vessel segments connect them.
- **Segment graph:** vessel segments are nodes; segments sharing an endpoint are connected.

Vessel Centerlines Become Graphs for pCR Prediction




Figure 2. Ultrafast breast MRI movie to vessel-network prediction.

A published deep-learning method (nnU-Net) identifies vessel voxels [2]. Our **tc4d** algorithm converts that mask into a one-voxel-wide centerline while preserving vascular connectivity. The centerline then becomes one of the three graph representations.

Nodes carry geometry and measured enhancement curves. A **graph neural network (GNN)** combines neighboring-node information, pools the network into one exam representation, and predicts pCR.

Forecasting Pretraining: A Controlled 2x2 Comparison




Figure 3. Forecasting weights transfer into the pCR classifier; a randomly initialized twin is the control arm.

The encoder receives vessel curves, connections, and elapsed times and predicts **future enhancement at connected vessel nodes** [4]. The downstream comparison holds inputs, encoder, training, patient-grouped folds, and evaluation fixed: forecasting versus random initialization, with and without graph message passing.

Centerline enhancement features predict pCR modestly, but hand-engineered GNNs do not beat matched tabular summaries—motivating this pretraining comparison as an alternative route to a better node representation.

Clinical value remains unproven: repeated seeds and locked-cohort validation must reproduce this gain before any clinical or vascular-specific claim.

Forecast Pretraining Improves the GNN in One Development Run




Figure 4. Patient-level pooled out-of-fold AUC with patient-bootstrap 95% CIs on shared patient-grouped folds.

Paired pretraining ΔAUC : **+0.149** [+0.046, +0.255] for the GNN, versus +0.034 [−0.049, +0.120] for graph-free (unresolved). The pretrained GNN exceeded pre-trained graph-free by +0.079 [+0.010, +0.155].

The claim is narrow: initialization helped this GNN on one seed, and its absolute interval includes chance.

Prediction Is Sensitive to Acquisition Timing

Acquisition timing is associated with outcome in this cohort: a model using frame count, duration, and cadence without imaging features reaches **0.587 AUC**—essentially the pretrained GNN's 0.583. Their predictions correlate (Pearson $r = 0.50$).

This raises concern about protocol dependence, but cannot separate technical from biological association. Adjustment can remove genuine patient differences that happen to track acquisition as well as technical variation.

Acknowledgments + References

We thank Zhen and Milica for feedback and MRI-physics expertise; Fred Howard for the development cohort; and Bella Summe, Julia Luo, José Cardona Arias, and Rebecca Wu for earlier vessel-pipeline development. HIRO (RRID:SCR_018372) coordinated imaging (Ludwig Fund; NIH/NCI P30 CA014599). Randi HPC/CRI compute was supported by BSD and ITM/CTSA NIH UL1TR002389.

- [1] Cortazar et al. Pathological complete response and long-term clinical benefit in breast cancer: The CTNeoBC pooled analysis. *The Lancet*, 384:164–172, 2014.
- [2] Isensee et al. nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18:203–211, 2021.
- [3] Nagy et al. Why are tumour blood vessels abnormal and why is it important to know? *British Journal of Cancer*, 100:865–869, 2009.
- [4] Tang et al. Self-supervised graph neural networks for improved electroencephalographic seizure analysis. *Proceedings of the AAAI Conference on Artificial Intelligence*, 36:12289–12297, 2022.